

# B1 Optimization

---

4 Lectures

1 Examples Sheet

Michaelmas 2023

Prof V A Prisacariu

---

- **Lecture 1:** Local and global optima, unconstrained univariate and multivariate optimization, stationary points, steepest descent.
- **Lecture 2:** Newton and Newton like methods – Quasi-Newton, Gauss-Newton; the Nelder-Mead (amoeba) simplex algorithm.
- **Lecture 3:** Linear programming constrained optimization; the simplex algorithm, interior point methods; integer programming.
- **Lecture 4:** Convexity, robust cost functions, methods for non-convex functions – grid search, multiple coverings, branch and bound, simulated annealing, evolutionary optimization.
- Course based on previous B1 Optimization, by Prof A Zisserman.


# Textbooks

---

## Practical Optimization

Philip E. Gill, Walter Murray, and Margaret H. Wright

Covers unconstrained and constrained optimization. Very clear and comprehensive.



# Background reading and web resources


---

## Numerical Recipes in C (or C++) : The Art of Scientific Computing

William H. Press, Brian P. Flannery, Saul A. Teukolsky, William T. Vetterling

CUP 1992/2002

- Good chapter on optimization
- Available on line as pdf



# Lecture 1

---

## Topics covered in this lecture:

- Problem formulation;
- Local and global optima;
- Unconstrained univariate optimization;
- Unconstrained multivariate optimization for quadratic functions:
  - Stationary points;
  - Steepest descent.

# Introduction

---

Optimization is used to find the best or optimal solution to a problem.

## **Steps involved in formulating an optimization problem:**


- Conversion of the problem into a mathematical model that abstracts all the essential elements;
- Choosing a suitable optimization method for the problem;
- Obtaining the optimum solution.

# Example: airfoil/wing aerodynamic shape optimization

---

## Optimization problem:

- constraints: lift, area, chord;
- minimize drag;
- vary shape.



# Introduction: Problem specification

---

Suppose we have a **cost function** (or **objective function**):

$$f(\mathbf{x}): \mathbb{R}^n \rightarrow \mathbb{R}$$

Our aim is find the value of the **parameters**  $\mathbf{x}$  that minimize the function:

$$\mathbf{x}^* = \arg \min_{\mathbf{x}} f(\mathbf{x})$$

subject to the following **constraints**:

- equality:  $c_i(\mathbf{x}) = 0, i = 1, \dots, m_e.$
- inequality:  $c_i(\mathbf{x}) \geq 0, i = m_e + 1, \dots, m.$

We will start by focussing on **unconstrained** problems.


# Recall: One dimensional functions

---


A differentiable function has a **stationary point** when the derivative is zero:

$$\frac{df}{dx} = 0$$


The second derivative gives the **type** of stationary point.



$f''(x) < 0$   
maximum



$f''(x) = 0$   
inflection




$f''(x) > 0$   
minimum



# Unconstrained optimization

---




- down-hill search (gradient descent) algorithms can find local minima;
- which of the minima is found depends on the starting point;
- such minima often occur in real applications.


## Example: template matching in 2D images

---

Model,  $M$



Transformation,  $T$



Data,  $D$

Input:

Two point sets  $M = \{\mathbf{M}_i\}$  and  $D = \{\mathbf{D}_j\}$ .

Task:

Determine the transformation  $T$  that minimizes the error between  $D$  and the transformed  $M$ .


Motivation:

Robot picking up “C”

# Cost function (correspondences known)

---

2D points  $(x, y)^T$ , Model  $\mathbf{M}_i$ , Data  $\mathbf{D}_i$



$$f(\theta, t_x, t_y) = \sum_i |\mathbf{R}(\theta)\mathbf{M}_i + \mathbf{t} - \mathbf{D}_i|^2$$


Transformation parameters:

- Rotation angle  $\theta$
- Translation  $\mathbf{t} = (t_x, t_y)^T$

# Cost function (correspondences unknown)

---

2D points  $(x, y)^T$ , Model  $\mathbf{M}_j$ , Data  $\mathbf{D}_i$



$$f(\theta, t_x, t_y) = \sum_i \min_j |\mathbf{R}(\theta)\mathbf{M}_j + t - \mathbf{D}_i|^2$$

for each  
data point

find closest model point

Transformation parameters:

- Rotation angle  $\theta$
- Translation  $\mathbf{t} = (t_x, t_y)^T$

# Cost function (matches unknown)


$$f(\theta, t_x, t_y) = \sum_i \underbrace{\min_j}_{\text{for each data point}} \underbrace{|\mathbf{R}(\theta)\mathbf{M}_j + \mathbf{t} - \mathbf{D}_i|}_{\text{find closest model point}}^2$$

As distances become smaller we get the correct correspondences and model pulled onto the data


Model point:  $\mathbf{M}_j = (x_j, y_j)^T$

Transformation parameters:

- Rotation angle  $\theta$
- Translation  $\mathbf{t} = (t_x, t_y)^T$




correct correspondences



closest point correspondences

# Performance


---



# Unconstrained **univariate** optimization

---

For the moment, assume we can start close to the global minimum



We will look at three basic methods to determine the minimum:

1. Gradient descent;
2. Polynomial interpolation;
3. Newton's method.



These introduce the ideas that will be applied in the **multivariate** case.

# A typical 1D function

---

As an example, consider the function:

$$f(x) = 0.1 + 0.1x + x^2 / (0.1 + x^2)$$






# 1. Gradient descent

---

Given a starting location,  $x_0$ , examine  $\frac{df}{dx}$  and move in the *downhill* direction to generate a new estimate  $x_1 = x_0 + \delta x$ .




$$\delta x = -\alpha \frac{df}{dx}$$

How to determine the step size  $\delta x$  ?

# Setting alpha ..

---



If the step size is too large, gradient descent may not converge (left)

If it is too small, convergence will be slow (middle).


A good step size leads to fast convergence (right)


## 2. Polynomial interpolation (trust region method)

---

Approximate  $f(x)$  with a simpler function which reasonably approximates the function in a neighbourhood around the current estimate  $x$ . This neighbourhood is the **trust region**.

- Bracket the minimum.
- Fit a quadratic or cubic polynomial which interpolates  $f(x)$  at some points in the interval.
- Jump to the (easily obtained) minimum for the polynomial.
- Throw away the worst point and repeat the process.





Quadratic interpolation using 3 points, 2 iterations

Other methods to interpolate a quadratic ?

- e.g. 2 points and one gradient

### 3. Newton's method

---

Fit a quadratic approx. to  $f(x)$  using both gradient and curvature information at  $x$ .

- Expand  $f(x)$  locally using a Taylor series:

$$f(x + \delta x) = f(x) + \delta x f'(x) + \frac{\delta x^2}{2} f''(x) + \text{h.o.t.}$$

- Find the  $\delta x$  (the variable here) such that  $x + \delta x$  is a stationary point of  $f$ :

$$\frac{d}{d\delta x} \left( f(x) + \delta x f'(x) + \frac{\delta x^2}{2} f''(x) \right) = f'(x) + \delta x f''(x) = 0$$


- and rearranging:

$$\delta x = -\frac{f'(x)}{f''(x)}$$


- Update for  $x$ :

$$x_{n+1} = x_n - \frac{f'(x_n)}{f''(x_n)}$$

## Newton iterations




## Quadratic approximations




- avoids the need to bracket the root;
- quadratic convergence (decimal accuracy doubles at every iteration).

## Newton iterations



## Quadratic approximations




- global convergence of Newton's method is poor;
- often fails if the starting point is too far from the minimum;
- in practice, must be used with a globalization strategy which reduces the step length until function decrease is assured.

# Stationary Points for Multidimensional functions

---

$f(x): \mathbb{R}^n \rightarrow \mathbb{R}$  has a stationary point with the gradient

$$\nabla f = \left( \frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_n} \right)^T = 0$$



$$\nabla f = \left( \frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2} \right)^T = 0$$



## Extension to N dimensions

---

- How big can N be?: problem sizes can vary from a handful of parameters to millions.
- In the following we will first **examine the properties of stationary points** in N dimensions and then move onto **optimization algorithms** to find the stationary point (minimum).
- We will consider examples for  $N = 2$ , so that cost function surfaces can be visualized.

# Taylor expansion in 2D

---

A function can be approximated locally by its Taylor series expansion about a point  $x_0$ .

$$f(x_0 + x) \approx f(x_0) + \left( \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right) \begin{pmatrix} x \\ y \end{pmatrix} + \frac{1}{2} (x, y) \begin{bmatrix} \frac{\partial^2 f}{\partial x^2} & \frac{\partial^2 f}{\partial x \partial y} \\ \frac{\partial^2 f}{\partial x \partial y} & \frac{\partial^2 f}{\partial y^2} \end{bmatrix} \begin{pmatrix} x \\ y \end{pmatrix} + \text{h.o.t.}$$

- This is a generalization of the 1D Taylor series:

$$f(x + \delta x) = f(x) + \delta x f'(x) + \frac{\delta x^2}{2} f''(x) + \text{h.o.t.}$$

- The expansion to second order is a **quadratic** function in  $\mathbf{x}$ .

$$f(\mathbf{x}) = a + \mathbf{g}^T \mathbf{x} + \frac{1}{2} \mathbf{x}^T \mathbf{H} \mathbf{x}$$

# Taylor expansion in ND

---

A function may be approximated locally by its Taylor expansion about a point  $x_0$

$$f(x_0 + x) \approx f(x_0) + \nabla f^T \mathbf{x} + \frac{1}{2} \mathbf{x}^T \mathbf{H} \mathbf{x} + \text{h.o.t.}$$

where the **gradient**  $\nabla f(\mathbf{x})$  of  $f(\mathbf{x})$  is the vector

$$\nabla f(\mathbf{x}) = \left[ \frac{\partial f}{\partial x_1}, \dots, \frac{\partial f}{\partial x_N} \right]^T$$

and the **Hessian**  $\mathbf{H}(\mathbf{x})$  of  $f(\mathbf{x})$  is the symmetric matrix

$$\begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_N} \\ \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_1 \partial x_N} & \cdots & \frac{\partial^2 f}{\partial x_N^2} \end{bmatrix}$$

# Properties of Quadratic functions

---

Taylor expansion:

$$f(\mathbf{x}_0 + \mathbf{x}) = f(\mathbf{x}_0) + \mathbf{g}^T \mathbf{x} + \frac{1}{2} \mathbf{x}^T \mathbf{H} \mathbf{x}$$

Expand about a stationary point  $\mathbf{x}_0 = \mathbf{x}^*$  in direction  $\mathbf{p}$ :

$$f(\mathbf{x}^* + \alpha \mathbf{p}) = f(\mathbf{x}^*) + \mathbf{g}^T \alpha \mathbf{p} + \frac{1}{2} \alpha^2 \mathbf{p}^T \mathbf{H} \mathbf{p} = f(\mathbf{x}^*) + \frac{1}{2} \alpha^2 \mathbf{p}^T \mathbf{H} \mathbf{p}$$

since the stationary point  $\mathbf{g} = \nabla f|_{\mathbf{x}^*} = 0$ .

At a stationary point the behavior is determined by  $\mathbf{H}$ .

# Properties of Quadratic functions

---

H is a symmetric matrix, so it has orthogonal eigenvectors

$$H\mathbf{u}_i = \lambda_i \mathbf{u}_i \text{ choose } |\mathbf{u}_i| = 1$$

$$f(\mathbf{x}^* + \alpha \mathbf{u}_i) = f(\mathbf{x}^*) + \frac{1}{2} \alpha^2 \mathbf{u}_i^T H \mathbf{u}_i = f(\mathbf{x}^*) + \frac{1}{2} \alpha^2 \lambda_i$$

As  $|\alpha|$  increases,  $f(\mathbf{x}^* + \alpha \mathbf{u}_i)$  increases, decreases or is unchanging according to whether  $\lambda_i$  is positive, negative or zero.

# Examples of Quadratic functions


---

Case 1: both eigenvalues positive


$$f(\mathbf{x}) = a + \mathbf{g}^T \mathbf{x} + \frac{1}{2} \mathbf{x}^T \mathbf{H} \mathbf{x}$$

with

$$a = 0, \mathbf{g} = \begin{bmatrix} -50 \\ -50 \end{bmatrix}, \mathbf{H} = \begin{bmatrix} 6 & 4 \\ 4 & 6 \end{bmatrix}$$



minimum



# Examples of Quadratic functions



---

## Case 2: eigenvalues have different signs

$$f(\mathbf{x}) = a + \mathbf{g}^T \mathbf{x} + \frac{1}{2} \mathbf{x}^T \mathbf{H} \mathbf{x}$$

with

$$a = 0, \mathbf{g} = \begin{bmatrix} -30 \\ +20 \end{bmatrix}, \mathbf{H} = \begin{bmatrix} 6 & 0 \\ 0 & -4 \end{bmatrix}$$



saddle surface: extremum but not a minimum

# Examples of Quadratic functions

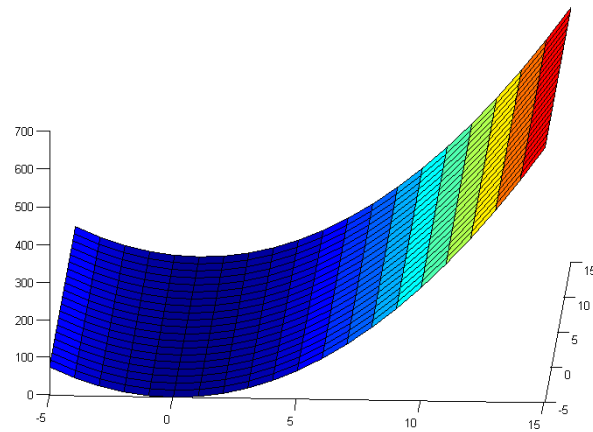
---

Case 3: one eigenvalue zero.


$$f(\mathbf{x}) = a + \mathbf{g}^T \mathbf{x} + \frac{1}{2} \mathbf{x}^T \mathbf{H} \mathbf{x}$$

with


$$a = 0, \mathbf{g} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}, \mathbf{H} = \begin{bmatrix} 6 & 0 \\ 0 & 0 \end{bmatrix}$$




parabolic cylinder








Hessian positive definite.  
Convex function.  
Minimum point.



Hessian negative definite.  
Concave function.  
Maximum point.





Hessian mixed.  
Surface has negative point.  
Saddle point.

# Optimization in N dimensions

---

- Reduce optimization in N dimensions to a series of (1D) line minimizations.
- Use methods developed in 1D (e.g. polynomial interpolation).



# Optimization in N dimensions

---

Start at  $\mathbf{x}_0$  then repeat:

1. Compute a search direction  $\mathbf{p}_k$ .
2. Compute a step length  $\alpha_k$ , such that  $f(\mathbf{x}_k + \alpha_k \mathbf{p}_k) < f(\mathbf{x}_k)$ .
3. Update  $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k$
4. Check for convergence (termination criteria), e.g.  $\nabla f \approx 0$ .

Reduce optimization in N dimensions to a series of (1D) line minimizations.

# Steepest descent

---

Basic principle is to minimize the N-dimensional function by a series of 1D line-minimizations:

$$\mathbf{x}_{n+1} = \mathbf{x}_n + \alpha_n \mathbf{p}_n$$

The steepest decent method choses  $\mathbf{p}_n$  to be parallel to the negative gradient:

$$\mathbf{p}_n = -\nabla f(\mathbf{x}_n)$$

The step-size  $\alpha_n$  is chosen to minimize  $f(\mathbf{x}_n + \alpha_n \mathbf{p}_n)$ .

For quadratic forms there is a closed form solution:


$$\alpha_n = -\frac{\mathbf{p}_n^T \mathbf{p}_n}{\mathbf{p}_n^T \mathbf{H} \mathbf{p}_n}$$

[exercise – minimize  $f(\mathbf{x}_n + \alpha_n \mathbf{p}_n)$  wrt  $\alpha$ ]

# Steepest descent example

---

$$a = 0, \mathbf{g} = \begin{bmatrix} -50 \\ -50 \end{bmatrix}, \mathbf{H} = \begin{bmatrix} 6 & 4 \\ 4 & 6 \end{bmatrix}$$



Steepest descent at  $\mathbf{x}_0 = [0, 14]$

After each line minimization, the new gradient is always **orthogonal** to the previous direction step

- True for any line minimization.
- Can be proven by examining the derivation for  $\alpha_n$ .

Consequently, the iterations may zig-zag down the valley in a very inefficient matter.

# Conjugate gradients

---

The method of conjugate gradients chooses successive descent directions  $\mathbf{p}_n$  such that it is guaranteed to reach the minimum in a finite number of steps.

- Each  $\mathbf{p}_n$  is chosen to be conjugate to all previous search directions wrt the Hessian  $H$ :


$$\mathbf{p}_n^T H \mathbf{p}_n = 0$$

- The resulting search directions are mutually linearly independent.
- Remarkably,  $\mathbf{p}_n$  can be chose using only knowledge of  $\mathbf{p}_{n-1}$ ,  $\nabla f(\mathbf{x}_{n-1})$  and  $\nabla f(\mathbf{x}_n)$  (see Numerical Recipes), e.g.

$$\mathbf{p}_n = \nabla f_n + \left( \frac{\nabla f_n^T \nabla f_n}{\nabla f_{n-1}^T \nabla f_{n-1}} \right) \mathbf{p}_{n-1}$$

# Conjugate gradients

---



- An N-dimensional quadratic form can be minimized in **at most N** conjugate descent steps.
- In figure: 3 different starting steps, minimum is reached in exactly 2 steps.

# What is next?

---

- Move from functions that are exactly quadratic to general functions that are represented locally by a quadratic
- Newton's method (that uses 2<sup>nd</sup> derivatives) and Newton-like methods for general functions